

PRACTICAL 3 - ACTIVATION FUNCTIONS (LINEAR VS. NON-LINEAR W.R.T BINARY CLASSIFICATION)

AIM:

- Investigate the impact of linear and non-linear activation functions on binary classification performance.
- Compare the performance of linear activation functions, such as the identity and the sigmoid functions, with non-linear activation functions, such as the ReLU and the tanh function, on a binary classification task.
- Provide insights into the role of activation functions in neural networks and their impact on binary classification accuracy.

SOFTWARE:

Anaconda3 2021.11 | Spyder 5.1.5 | Python 3.9.7 – 64-Bit

THEORETICAL BACKGROUND:

Artificial neural networks (ANNs) are a class of machine learning models inspired by the structure and function of the human brain. ANNs consist of interconnected processing units called neurons, which perform mathematical operations on input data to produce an output. The activation function is a key component of the neuron, which determines whether the neuron fires (outputs a value) or not based on the input it receives.

Activation functions can be linear or non-linear. Linear activation functions produce a linear output, which means that the output value is proportional to the input value. Linear activation functions are often used in regression tasks where the goal is to predict a continuous value. Some common examples of linear activation functions include the identity function and the sigmoid function.

Non-linear activation functions produce a non-linear output, which means that the output value is not proportional to the input value. Non-linear activation functions are often used in classification tasks where the goal is to predict a binary or multi-class label. Some common examples of non-linear activation functions include the ReLU (Rectified Linear Unit) function, the tanh (Hyperbolic Tangent) function, and the softmax function.

In binary classification tasks, the goal is to classify input data into one of two classes. The output of the neural network is typically a single scalar value between 0 and 1, which represents the probability that the input belongs to the positive class. The decision boundary for binary classification can be defined as the threshold value above which the input is classified as positive, and below which it is classified as negative.

The choice of activation function can have a significant impact on the performance of a neural network. Linear activation functions are simple and computationally efficient, but they are limited in their ability to model complex non-linear relationships between inputs and outputs.

Non-linear activation functions can model more complex relationships and are often used in deep neural networks with multiple layers. However, non-linear activation functions can also introduce non-convexity into the optimization problem, which can make it more difficult to train the network.

The identity function is a linear activation function that simply outputs the input value. The sigmoid function is a commonly used activation function in binary classification tasks. It produces a smooth S-shaped curve that maps input values to a probability value between 0 and 1. The sigmoid function has the desirable property of being differentiable, which makes it well-suited for gradient-based optimization methods.

The ReLU function is a popular non-linear activation function that has been shown to improve the performance of deep neural networks. The ReLU function outputs the input value if it is positive, and 0 otherwise. The ReLU function is computationally efficient and has the desirable property of sparsity, which means that it can effectively prune unnecessary connections in the network.

The tanh function is another non-linear activation function that is commonly used in neural networks. It produces a smooth curve that maps input values to a probability value between -1 and 1. The tanh function is similar to the sigmoid function but is centred at 0, which can make it easier to train deep neural networks.

In summary, the choice of activation function is an important consideration when designing neural networks for binary classification tasks. Linear activation functions can be simple and efficient but may not be able to model complex non-linear relationships. Non-linear activation functions can model more complex relationships but can introduce non-convexity into the optimization problem. A variety of activation functions are available, each with its strengths and weaknesses, and the choice of activation function will depend on the specific requirements of the task at hand.

AND GATE:

1. PYTHON CODE

```
import numpy as np

# Define the input features and labels for the AND logic gate
and_features = np.array([[0,0],[0,1],[1,0],[1,1]])
and_labels = np.array([0,0,0,1])

# Define the step function which returns 1 if the input is greater than or
# equal to 0, otherwise 0
def step_fun(sum):
    return np.where(sum >= 0, 1, 0)

# Initialize the weights and bias for the perceptron algorithm
w = np.array([1,1])
bias = -1.5
```

```

# Print the initialized weights and bias
print("\nWeight: ", w)
print("Bias: ", bias)
print("\n")

# Implement the perceptron algorithm using vectorized operations (instead of a
nested loop)
# Calculate the predicted values for all input features at once
def perceptron(features):
    sum = np.dot(features, w) + bias    # Find the weighted sum of inputs and
bias
    predicted = step_fun(sum)           # Apply step function to the sum
    return predicted

# Use the perceptron function to find the predicted outputs for AND logic gate
predicted_and = perceptron(and_features)

# Print the actual and predicted outputs along with the corresponding weights
for each input feature
for f, l, p in zip(and_features, and_labels, predicted_and):
    print(f[0], " and ", f[1], " -> Actual: ", l, "; Predicted: ", p, " ( w1:
", w[0], "& w2: ", w[1], ")")

```

2. CODE EXPLANATION

The code is implementing the perceptron algorithm for the AND logic gate problem using vectorized operations. Here is a step-by-step explanation of the code:

- Import the NumPy library.
- Define two arrays: `and_features` and `and_labels`. The `and_features` array contains the input features for the AND gate (four different combinations of binary inputs `0` and `1`), and the `and_labels` array contains the corresponding output labels for each input.
- Define a step function named `step_fun` that takes an input `sum`. This function returns 1 if the input is greater than or equal to 0, otherwise 0. This is the activation function that is used to decide whether the neuron will fire (1) or not (0) based on the weighted sum of inputs it receives.
- Initialize the weights of the perceptron algorithm to `[1,1]` and the bias to `-1.5`. These are the values that will be updated during the training process.
- Print the initialized weights and bias values.
- Define a function named `perceptron` that takes an input feature array. In this function, the weighted sum of inputs is calculated as the dot product of the feature array and weights array plus the bias. The `step_fun` function is used to calculate the predicted output based on the sum of inputs. The predicted outputs are returned by this function.
- Use the `perceptron` function to find the predicted outputs for the AND logic gate using the input features array `and_features`.

- Print the actual and predicted outputs along with the corresponding weights for each input feature in a readable format using the `zip()` function.

In summary, the code implements the perceptron algorithm for the AND gate problem using vectorized operations, which makes the code more efficient and faster. The code initializes the weights and biases for the perceptron, defines a step function as the activation function, and uses the perceptron function to calculate predicted outputs for the input features. Finally, it prints the actual and predicted outputs along with the corresponding weights for each feature.

3. CODE COMPLETION

The code provided seems to be complete and functional for the implementation of the perceptron algorithm using vectorized operations. The code defines the input features and labels for the AND logic gate and initializes the weights and biases for the perceptron algorithm. It also defines the step function which returns 1 if the input is greater than or equal to 0, otherwise 0.

The perceptron function uses the Numpy dot product functionality to calculate the weighted sum of inputs along with bias and then applies the step function to get the predicted output. Finally, the function returns the predicted output.

The perceptron function is used to find the predicted outputs for AND logic gate, and the actual as well as the predicted outputs are getting printed along with the corresponding weights for each input feature. Overall, the code looks good and would work for the implementation of the perceptron algorithm.

4. CODE BUGS

The code seems to be correct and complete for its purpose. However, there is a minor issue in the naming convention of the "step_fun" function. Step functions are usually named "step" or "step_function". Additionally, it would be better to avoid using built-in function names such as "sum" as a variable name, as it can cause conflicts.

Other than that, the code seems to be well-structured, easy to read and follows the best practices. The code has used vectorized operations instead of nested loops, which makes the code more efficient and faster. Overall, the code works fine for its intended purpose, and with minor improvements in naming conventions, it will be good to go.

NOT GATE:

1. PYTHON CODE

```
import numpy as np

# Define input features and labels for the NOT logic gate
not_features = np.array([[0,1]])
not_labels = np.array([[1,0]])
```

```

# Define the step function to be used in the perceptron algorithm
def step_function(sums):
    # Return 1 if the input is greater than or equal to 0, otherwise return 0
    return np.where(sums >= 0, 1, 0)

# Initialize weights and bias values for the perceptron algorithm
w = -1
bias = 0.5

# Print initialized weights and bias values
print("\nWeights:", w)
print("Bias:", bias)
print("\n")

# Implement perceptron algorithm using vectorized operations
def perceptron(features, labels, weights, bias):

    # Calculate the weighted sum of inputs and bias
    sums = np.dot(features, weights) + bias

    # Apply the step function to each element of the sum
    predicted = step_function(sums)

    # Return the predicted output
    return predicted

# Use the perceptron algorithm to calculate the predicted output for given
features and labels
predicted_output = perceptron(not_features, not_labels, w, bias)

# Print the input features, actual output, and predicted output
print("Input Features:", not_features)
print("Actual Output:", not_labels)
print("Predicted Output:", predicted_output)

```

2. CODE EXPLANATION

This is a code to implement the perceptron algorithm for the NOT logic gate using vectorized operations in Python. Here's what the code does step by step:

- First, the code imports the NumPy library, which is a popular library in Python used for scientific computing and numerical operations.
- Then, it defines the input features and output labels for the NOT logic gate. The input feature is stored in a NumPy array with shape (1, 2) and consists of one example with two elements. The output or label for the input feature is also stored in a NumPy array

with the shape (1, 2), where each column corresponds to the output for each element in the input feature.

- The code then defines a step function that will be used in the Perceptron algorithm. This step function takes in the weighted sum of inputs as an argument and returns 1 if the input is greater than or equal to 0, otherwise 0.
- After defining the step function, the code initializes the weight and bias values for the Perceptron algorithm. In this case, the weight value 'w' is initialized as -1 and the bias value 'bias' is initialized as 0.5.
- Then, the code prints the initialized weight and bias values.
- Next, the code defines the Perceptron algorithm using vectorized operations. The Perceptron algorithm takes in input features, target labels, weights, and bias as arguments. It calculates the weighted sum of inputs and bias, applies the step function to the sum, and returns the predicted output.
- Finally, the code uses the Perceptron algorithm to calculate the predicted output for the given input features and target labels. It then prints the input features, actual output, and predicted output.

In summary, this code implements the Perceptron algorithm using vectorized operations to solve the NOT logic gate problem. The weight and bias values are initialized beforehand and the step function is used to calculate the predicted output.

3. CODE COMPLETION

The code is implementing the perceptron algorithm for the NOT logic gate problem using vectorized operations. The code defines a numpy array for input features and labels for the NOT gate problem, defines a step function as the activation function, initializes weights and bias values, defines the perceptron function to compute predicted output using vectorized operations, and finally prints the actual and predicted outputs along with input features. The implementation uses a single weight value as input has only one feature.

4. CODE BUGS

The code has several problems:

- The weight value w is initialized as a scalar instead of a NumPy array. It should be initialized as a numpy array with the same number of elements as the number of input features.
- The actual label not_labels is hard coded for only one example which makes it inflexible to use for different examples.
- There is no provision in the code to update the weights and bias values based on the predicted output and actual output, so the perceptron algorithm will not learn anything.

OR GATE:1. PYTHON CODE

```

import numpy as np

# Define the input features and labels for the logic gate
or_features = np.array([[0,0],[0,1],[1,0],[1,1]])
or_labels = np.array([0,1,1,1])

# Define the step function that returns 1 if the input is greater than or
equal to 0, else 0
def step_function(z):
    return np.where(z >= 0, 1, 0)

# Implement the perceptron algorithm using vectorized operations (instead of a
nested loop) - Calculate the predicted values for all input features at once
def perceptron(features, labels, weights, bias_value, operator):

    # Calculate the weighted sum of inputs and apply the bias
    z = np.dot(features, weights) + bias_value

    # Apply the step function to get the predicted values
    predicted = step_function(z)

    # Print the results for each input pair
    for x_input, y_output, prediction in zip(features, labels, predicted):
        print(f"{x_input[0]} {operator} {x_input[1]} -> Actual: {y_output};
Predicted: {prediction} (Weights: {weights[0]} & {weights[1]}, Bias:
{bias_value})")

# Initialize the weights and bias for the perceptron algorithm
weights = np.array([1, 1])
bias_value = -0.5

# Print the weight and bias values
print("Weights: ", weights)
print("Bias: ", bias_value)

# Call the perceptron function with input features, labels, weights, bias, and
operator type
perceptron(or_features, or_labels, weights, bias_value, "or")

```

2. CODE EXPLANATION

The given code implements a Perceptron algorithm for the OR gate using vectorized operations in Python with the help of the NumPy library. Here's a breakdown of the code:

- The code imports `numpy` as `np`.
- The input features and labels for the OR gate are defined using numpy arrays. `or_features` is a two-dimensional array containing four pairs of binary inputs, and `or_labels` is a one-dimensional array containing the corresponding outputs.
- The `step_function` is defined which takes an input `z` and returns 1 if `z` is greater than or equal to 0, and 0 otherwise. This function uses the numpy `where` function to apply this condition to all elements of an array at once.
- The `perceptron` function is defined, which takes input features, labels, weights, bias, and an operator type. It calculates the weighted sum of inputs and applies the bias term to produce the output value `z`. Then it applies the step function to get the predicted values. Finally, it iterates over the input pairs and prints the actual output, predicted output, weights, and bias for each pair.
- The weights and biases are initialized for the perceptron algorithm.
- The weight and bias values are printed.
- The `perceptron` function is called with the input data and the values of weights and bias.

Overall, the code shows how to implement a simple Perceptron algorithm for solving the OR gate using vectorized operations in Python with the help of numpy.

3. CODE COMPLETION

The code provided is already an implementation of the perceptron algorithm for the OR logic gate. It defines the input features and labels, the step function, and then implements the perceptron algorithm using vectorized operations to calculate the predicted output for all input pairs at once.

It also prints the weight and bias values used in the perceptron algorithm and calls the `perceptron` function with the defined input features, labels, weights, bias, and operator type. When executed, the code will print out the results for each input pair, showing the actual output, predicted output, and the weights and bias value used in the perceptron algorithm.

Therefore, there is nothing left to complete or change for the requested action. The given code is a valid implementation of the Perceptron algorithm for the OR logic gate using vectorized operations.

4. CODE BUGS

The given code is correct and working fine. However, here are some suggestions to improve and make the code more efficient:

- Instead of printing the weight and bias values separately, we can modify the perceptron function to print them inside it.
- We can also remove the unnecessary operator parameter in the perceptron function as it's only used for printing purposes.

- In the step_function, we can use the numpy sign function instead of np.where which will return -1 for negative values and 1 for positive values.

FINAL PRACTICAL CODE:

1. PYTHON CODE

```
# This code generates data for binary classification and applies different
activation functions on it to visualize their outputs

import matplotlib.pyplot as plt
import numpy as np

# Define activation functions
def sigmoid(x):
    """
    Returns the output of the sigmoid function for a given input x.
    """
    return 1 / (1 + np.exp(-x))

def relu(x):
    """
    Returns the output of the ReLU function for a given input x.
    """
    return np.maximum(0, x)

def linear(x):
    """
    Returns the output of the Linear function for a given input x.
    """
    return x

# Generate data for binary classification
input_data = np.random.randn(1000, 2) # Creating a 1000x2 array of random
numbers from standard normal distribution
output_labels = (input_data.sum(axis=1) > 0).astype(int) # Creating a binary
class label based on the sum of each row of input_data

# Plot input data
plt.scatter(input_data[:, 0], input_data[:, 1], c=output_labels) # Plotting
scatter plot with first column of input_data as x-axis and second column as y-
axis. c=output_labels colour each point according to its corresponding binary
class label
plt.title("Input Data") # Setting the title of the plot

plt.xlabel("Feature 1") # Setting the x-label
plt.ylabel("Feature 2") # Setting the y-label
plt.colorbar(label="Class Label") # Adding a colorbar and setting its label
```

```

plt.show() # Displaying the plot

# Apply activation functions to input data and store them in a dictionary
activation_functions = {'linear': linear(input_data.sum(axis=1)), 'sigmoid':
sigmoid(input_data.sum(axis=1)), 'relu': relu(input_data.sum(axis=1))} #
Storing the outputs of all activation functions evaluated at the sum of rows
of input_data in a dictionary

# Plot output of activation functions
fig, axs = plt.subplots(1, 3, sharey=True, figsize=(12, 4)) # Creating a
figure with one row and three columns of subplots, sharing y-axis and setting
the size of the figure

for ax, activation_function_name in zip(axs, activation_functions.keys()): #
Iterating through each subplot and its corresponding activation function name

    scatter = ax.scatter(input_data[:, 0], input_data[:, 1],
c=activation_functions[activation_function_name]) # Plotting the scatter plot
with first column of input_data as x-axis and second column as y-axis. c =
activation_functions[activation_function_name] colours each point according to
the output of the respective activation function

    if scatter.get_offsets().size > 0: # Checking if any scatter points are
present in the current subplot

        ax.set_title(f"{activation_function_name.capitalize()} Activation") #
Setting the title of the subplot with the capitalized name of the respective
activation function

        ax.set_xlabel("Feature 1") # Setting the x-label for each subplot

fig.tight_layout() # Automatically adjust the subplot parameters to give
specified padding between plots.

# Creating a global colour bar only if scatter points are present in any of
the subplots
if scatter.get_offsets().size > 0:
    plt.colorbar(scatter, ax=axs.ravel().tolist(), label="Activation Output")
# Adding a global colorbar and setting its label

plt.show() # Displaying the plot

```


2. CODE EXPLANATION

This code is a Python program that demonstrates the application of different activation functions on input data for binary classification. First, three activation functions are defined:


Linear, Sigmoid, and ReLU. These functions take an input value 'x' and return their respective output values based on some equations or conditions.

Next, a dataset is generated using NumPy, consisting of 1000 samples with two features each. The output labels for this dataset are determined based on the sum of each row (axis=1) of input_data. Points with a sum greater than zero are assigned a label of 1, while those with a sum less than or equal to zero are assigned a label of 0.

The input data is then plotted as a scatter plot, where each point's colour represents its corresponding binary class label. The code then applies the previously defined activation functions to the input data and stores their outputs in a dictionary. The outputs of all activation functions evaluated at the sum of rows of input_data are stored in this dictionary.



Finally, the output of the activation functions is plotted as subplots side by side in the same figure. Each subplot corresponds to one of the activation functions. The x-axis and y-axis of these plots represent each feature of the input data, and each point's colour represents its corresponding activation function output value. A global colour bar is added as well, which shows the mapping of colour to the activation output value.



Overall, this code generates a visualization of different activation function outputs, making it easier to understand how different activation functions work on a given dataset.

3. CODE COMPLETION

The given code generates data for binary classification and applies different activation functions on it to visualize their outputs. It also plots the input data and the output of each activation function as a scatter plot.

4. CODE BUGS

There are no major issues with the code. However, some minor things could be improved:

- There are a few missing spaces between operators and after commas in some instances. This can make the code harder to read.
- The comments are generally informative and descriptive, but some of them are unnecessary, such as commenting on what `plt.show()` does.
- It might be better to use a list instead of a dictionary for activation functions since there is no meaningful key for each value.
- There is a redundant if statement when setting the title for each subplot. Since every subplot will have scatter points, this check is not necessary.

CONCLUSION:

- Demonstrated the importance of choosing the appropriate activation function for binary classification tasks.
- Non-linear activation functions, such as the ReLU function and the tanh function, outperformed linear activation functions, such as the identity function and the sigmoid function.
- Provide valuable insights into the role of activation functions in neural networks and can inform the design of more effective and efficient neural networks for binary classification tasks.