

## **PRACTICAL 8 – IMPLEMENTING A RESTRICTED BOLTZMANN MACHINE (RBM) AND USE IT FOR DIMENSIONALITY REDUCTION**

### **AIM:**

- Implement a Restricted Boltzmann Machine (RBM) and apply it to perform dimensionality reduction on a given dataset.
- Effectively reduce the high-dimensional input data into a lower-dimensional representation, while retaining as much relevant information as possible.
- Demonstrate the RBM's ability to capture meaningful features and patterns in the data, providing a compressed representation that can facilitate improved data visualization, and storage, and potentially enhance downstream machine-learning tasks.

### **SOFTWARE:**

Anaconda3 2021.11 | Spyder 5.1.5 | Python 3.9.7 – 64-Bit

### **THEORETICAL BACKGROUND:**

#### **Dimensionality Reduction:**

Dimensionality reduction is a fundamental technique in machine learning and data analysis that aims to reduce the number of features or variables in a dataset while preserving the essential information. High-dimensional data can be computationally expensive and may suffer from the curse of dimensionality, where the performance of machine learning algorithms deteriorates as the number of features increases. Dimensionality reduction techniques help address these challenges by transforming the data into a lower-dimensional space, which can lead to improved model performance, easier visualization, and faster processing.

#### **Restricted Boltzmann Machine (RBM):**

The Restricted Boltzmann Machine (RBM) is a type of artificial neural network that belongs to the family of generative stochastic networks. Introduced by Paul Smolensky in the 1980s and popularized by Geoffrey Hinton in the early 2000s, RBMs are particularly well-suited for unsupervised learning tasks, such as collaborative filtering, feature learning, and, in this case, dimensionality reduction.

#### **RBM Structure:**

An RBM consists of two layers: the visible layer (input layer) and the hidden layer. Each layer contains a set of binary units, where the visible units correspond to the input data's features, and the hidden units learn to represent higher-level features or patterns in the data. These layers are fully connected, meaning each visible unit is connected to every hidden unit and vice versa.

Energy-Based Model:

RBM's are considered energy-based models, and they learn by minimizing an energy function associated with the configuration of visible and hidden units. The energy function is defined as the negative log-likelihood of the model's parameters and the data. The model aims to find the set of weights that minimizes the energy for the observed data while maximizing the energy for unobserved or corrupted data, thus effectively capturing the underlying probability distribution of the input data.

Training RBMs:

Training an RBM involves using a learning algorithm, such as contrastive divergence (CD) or persistent contrastive divergence (PCD), to adjust the weights to minimize the energy function. During training, the RBM learns to reconstruct the input data from a reduced set of hidden units, and it tries to model the probability distribution of the data effectively.

Dimensionality Reduction with RBMs:

To perform dimensionality reduction using an RBM, the model is trained on the input data with a higher number of features (high-dimensional data). After training, the model's weights are fixed, and the hidden layer activations are used as the reduced representation of the data. By using the hidden layer activations, which have a lower dimensionality compared to the input data, the RBM effectively encodes the most relevant and meaningful features of the original data in the reduced representation.

In this practical experiment, we will implement a Restricted Boltzmann Machine, train it on a given dataset, and utilize its learned hidden layer activations for dimensionality reduction. The experiment aims to assess the RBM's ability to capture essential features, effectively reduce the dimensionality of the data, and evaluate the impact of the reduced representation on downstream tasks, such as data visualization or classification.

**RAW CODE – AUTOENCODERS:****1. PYTHON CODE**

```
import numpy as np
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.preprocessing.image import ImageDataGenerator

# Enable GPU acceleration
physical_devices = tf.config.experimental.list_physical_devices('GPU')
if len(physical_devices) > 0:
    tf.config.experimental.set_memory_growth(physical_devices[0], True)
```

```

# Load the MNIST dataset
(x_train, y_train), (x_test, y_test) = keras.datasets.mnist.load_data()

# Normalize and reshape the input data
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0
x_train = np.reshape(x_train, (len(x_train), 28, 28, 1))
x_test = np.reshape(x_test, (len(x_test), 28, 28, 1))

# Apply data augmentation
datagen = ImageDataGenerator(
    rotation_range=10,
    zoom_range=0.1,
    width_shift_range=0.1,
    height_shift_range=0.1)

# Define the autoencoder model
input_img = keras.Input(shape=(28, 28, 1))
encoded = layers.Flatten()(input_img)
encoded = layers.Dense(64, activation='relu')(encoded)
decoded = layers.Dense(784, activation='sigmoid')(encoded)
decoded = layers.Reshape((28, 28, 1))(decoded)
autoencoder = keras.Model(input_img, decoded)
autoencoder.compile(optimizer='adam', loss='binary_crossentropy')

# Use early stopping to prevent overfitting
early_stopping = EarlyStopping(monitor='val_loss', min_delta=0.001,
patience=5)

# Train the autoencoder with augmented data and early stopping
autoencoder.fit(datagen.flow(x_train, x_train, batch_size=128),
    epochs=5,
    batch_size=128,
    shuffle=True,
    validation_data=(x_test, x_test),
    callbacks=[early_stopping])

# Test the autoencoder
decoded_imgs = autoencoder.predict(x_test)

# Display some results
import matplotlib.pyplot as plt
n = 10 # Number of digits to display
plt.figure(figsize=(20, 4))
for i in range(n):
    # Original images
    ax = plt.subplot(2, n, i + 1)

```

```
plt.imshow(x_test[i].reshape(28, 28))
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)
# Decoded images
ax = plt.subplot(2, n, i + 1 + n)
plt.imshow(decoded_imgs[i].reshape(28, 28))
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)
plt.show()
```

## 2. CODE EXPLANATION

The code is an example of an autoencoder model using the TensorFlow and Keras libraries in Python. Here is a breakdown of the code:

- *Importing Libraries:*
  - ``numpy`` (imported as ``np``): Used for numerical operations.
  - ``tensorflow`` (imported as ``tf``): A popular deep learning library.
  - ``keras``: A high-level neural networks API that runs on top of TensorFlow.
  - ``layers`` from ``keras``: Contains various types of layers used in building neural networks.
  - ``EarlyStopping`` from ``keras.callbacks``: Used to stop training if a monitored metric does not improve.
  - ``ImageDataGenerator`` from ``keras.preprocessing.image``: Used for data augmentation.
- *Enabling GPU Acceleration:*
  - The code checks if there are any physical GPUs available and sets memory growth for the first GPU (if present) to avoid hogging all the GPU memory.
- *Loading the MNIST Dataset:*
  - The code loads the MNIST dataset, which consists of handwritten digits and their corresponding labels.
  - It splits the dataset into training and test sets.
- *Normalizing and Reshaping the Input Data:*
  - The pixel values of the input images are normalized to the range [0, 1].
  - The shapes of the input arrays are reshaped to add a colour channel dimension.
- *Data Augmentation:*
  - An ``ImageDataGenerator`` object is created with specific augmentation parameters such as rotation, zoom, and shift ranges.
  - This object will be used to generate augmented images during training.
- *Defining the Autoencoder Model:*
  - The autoencoder model is defined using Keras' functional API.
  - The input layer takes in images with shapes (28, 28, 1).
  - The input is flattened, followed by a dense layer with ReLU activation to create an encoded representation.

- The encoded representation is then decoded back to the original image shape using dense and reshape layers.
- The autoencoder model is created using the ``Model`` class from Keras, with input and output defined.
- The model is compiled with the Adam optimizer and binary cross-entropy loss.
- *Early Stopping:*
  - An ``EarlyStopping`` callback is created to monitor the validation loss during training.
  - If the validation loss does not improve by at least 0.001 for 5 consecutive epochs, training will be stopped early.
- *Training the Autoencoder:*
  - The autoencoder model is trained using the ``fit`` method.
  - Augmented data is generated in batches using ``datagen.flow``.
  - Training runs for 5 epochs, with a batch size of 128 and shuffling of data.
  - Validation data is provided to monitor the performance during training.
  - The ``EarlyStopping`` callback is passed to prevent overfitting.
- *Testing the Autoencoder:*
  - The trained autoencoder is used to predict the reconstructed images for the test set.
- *Displaying Results:*
  - The original test images and their corresponding reconstructed images are displayed using Matplotlib.
  - A subplot layout is created to display multiple images.
  - The images are shown side-by-side, with the original images on the top row and the decoded images on the bottom row.

Overall, this code demonstrates how to build and train an autoencoder model using TensorFlow and Keras. It also shows how to use data augmentation and early-stopping techniques to improve performance and prevent overfitting.

### 3. CODE BUGS

The code does not have any major problems. However, there are a few things that can be improved or could potentially cause issues:

- *GPU Acceleration:* The code checks if there is at least one GPU available and then sets memory growth for the first GPU in the list. This assumes that there is only one GPU present and may not work correctly if multiple GPUs are available. To handle multiple GPUs, you can iterate over ``physical_devices`` and enable memory growth for each GPU.
- *Early Stopping:* The code uses early stopping with a ``min_delta`` of 0.001 and ``patience`` of 5. These values can vary depending on the problem and dataset. It's important to tune these hyperparameters based on the specific problem to get optimal results.
- *Displaying Images:* The code displays 10 original and decoded images using Matplotlib. It would be better to check if there are at least 10 images in the test set

before running the loop to avoid an index error. You can add a condition before the loop like this: `for i in range(min(n, len(x\_test))):`.

- *Plot Size:* The current size of the plot created by `plt.figure(figsize=(20, 4))` may not be suitable for all screen sizes. You might consider adjusting the figure size based on your needs and the size of the images being displayed.

Overall, the code looks good and should work fine with these improvements.

#### 4. CODE OPTIMIZATION

The code can be optimized in a few ways:

- *GPU Acceleration:* You are already enabling GPU acceleration by setting memory growth for the first GPU in the list. However, if there are multiple GPUs available, you can enable memory growth for all of them by iterating over `physical\_devices`.
- *Normalization and Reshaping:* Instead of using NumPy to normalize and reshape the input data, you can use the built-in functions provided by Keras.
- *Data Augmentation:* To save memory and improve efficiency, you can apply data augmentation during training with a smaller batch size. This can be achieved by applying the `datagen.flow` directly inside the `fit` method.
- *Displaying Images:* Since you are displaying only 10 images, you can avoid resizing the plot by using the `subplots` function with the `squeeze=False` argument. This allows you to directly access each subplot using the returned array of Axes objects.

By applying these optimizations, the code will be more efficient and easier to read.

### RAW CODE – AUTOENCODERS EXERCISE:

#### 1. PYTHON CODE

```
import tensorflow as tf
from tensorflow.keras.datasets import fashion_mnist
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, Dense
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint
import numpy as np
import matplotlib.pyplot as plt

# Load data
(train_images, train_labels), (test_images, test_labels) =
fashion_mnist.load_data()

# Normalize data
train_images = train_images.astype('float32') / 255.
test_images = test_images.astype('float32') / 255.

# Add random noise to the training images
```

```

noise_factor = 0.5
train_noisy = train_images + noise_factor * np.random.normal(loc=0.0,
scale=1.0, size=train_images.shape)
train_noisy = np.clip(train_noisy, 0., 1.)

# Add random noise to the test images
test_noisy = test_images + noise_factor * np.random.normal(loc=0.0, scale=1.0,
size=test_images.shape)
test_noisy = np.clip(test_noisy, 0., 1.)

# Reshape data for autoencoder model
train_images = train_images.reshape(train_images.shape[0], 784)
test_images = test_images.reshape(test_images.shape[0], 784)

# Define the autoencoder model
input_img = Input(shape=(784,))
encoded = Dense(128, activation='relu')(input_img)
encoded = Dense(64, activation='relu')(encoded)
encoded = Dense(32, activation='relu')(encoded)
decoded = Dense(64, activation='relu')(encoded)
decoded = Dense(128, activation='relu')(decoded)
decoded = Dense(784, activation='sigmoid')(decoded)

autoencoder = Model(input_img, decoded)

# Compile the model
autoencoder.compile(optimizer='adam', loss='binary_crossentropy')

# Set up callbacks
early_stopping = EarlyStopping(monitor='val_loss', patience=3, verbose=1,
mode='min')
checkpoint = ModelCheckpoint('best_model.h5', monitor='val_loss', verbose=1,
save_best_only=True, mode='min')

# Train the model
history = autoencoder.fit(
    train_noisy.reshape(train_noisy.shape[0], 784),
    train_images.reshape(train_images.shape[0], 784),
    epochs=50,
    batch_size=256,
    shuffle=True,
    validation_data=(
        test_noisy.reshape(test_noisy.shape[0], 784),
        test_images.reshape(test_images.shape[0], 784)
    ),
    callbacks=[early_stopping, checkpoint],
    verbose=1
)

```

```

# Load the best-saved model
autoencoder = tf.keras.models.load_model('best_model.h5')

# Evaluate the model on test data and display the loss
decoded_imgs = autoencoder.predict(test_noisy.reshape(test_noisy.shape[0],
784))
loss = autoencoder.evaluate(test_images.reshape(-1, 784), decoded_imgs)
print('Test loss:', loss)

# Display the original and reconstructed images
n = 10 # number of images to display
plt.figure(figsize=(20, 4))
for i in range(n):
    # display original images
    ax = plt.subplot(2, n, i + 1)
    plt.imshow(test_images[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)

    # Display reconstructed images
    ax = plt.subplot(2, n, i + 1 + n)
    plt.imshow(decoded_imgs[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
plt.show()

```

## 2. CODE EXPLANATION

This code implements an autoencoder model using TensorFlow and Keras. The autoencoder is trained on the Fashion MNIST dataset to reconstruct noisy images. The model consists of an encoder and a decoder, with progressively decreasing and increasing dimensions respectively.

The training process involves adding random noise to the input images and minimizing the difference between the reconstructed and original images. The best model is saved based on validation loss and used for evaluation. Finally, the code displays the original and reconstructed images using Matplotlib.

## 3. CODE COMPLETION

The code is an implementation of an autoencoder using TensorFlow and Keras. The autoencoder is trained to reconstruct images that have been corrupted with random noise. It is a complete script that loads the Fashion MNIST dataset, normalizes the data, adds random noise to the training and test images, defines the autoencoder model architecture, compiles the



model, sets up callbacks for early stopping and saving the best model, trains the model, evaluates it on the test data, and finally displays the original and reconstructed images.

#### 4. CODE BUGS

The code looks generally correct and should work as intended. However, there are a few potential issues that you may want to consider:

- *Data Preprocessing*: The code normalizes the pixel values of the images by dividing them by 255 to scale them between 0 and 1. This is a common preprocessing step for image data in machine learning tasks. However, it is important to ensure that both the training and test datasets are normalized using the same normalization factors.
- *Noise Level*: The variable ``noise_factor`` controls the amount of random noise added to the images. In the given code, ``noise_factor`` is set to 0.5. This value might be too high, resulting in noisy or distorted images. You can experiment with different values to find an optimal noise level.
- *Model Architecture*: The autoencoder model defined in the code consists of three encoding layers and three decoding layers. Each layer has a decreasing number of units until reaching the bottleneck layer (encoded representation). This architecture is a common choice for simple autoencoders, but it may not necessarily yield the best performance for your specific task. You could try experimenting with different architectures to improve the results.
- *Callbacks*: The code sets up two callbacks: ``EarlyStopping`` and ``ModelCheckpoint``. These callbacks are useful for monitoring the validation loss during training and saving the best model based on this criterion. However, the ``early_stopping`` callback is set to stop training if the validation loss does not decrease for 3 consecutive epochs. You can adjust the patience parameter according to your needs.
- *Displayed Images*: The code displays original and reconstructed images using matplotlib. The variable ``n`` determines the number of images to display, which is currently set to 10. You can modify this value to display more or fewer images, depending on your preferences.

Overall, these potential issues are largely subjective and depend on the specific requirements and goals of your project. It is recommended to experiment with different settings and architectures to achieve the desired results.

#### 5. CODE OPTIMIZATION

To optimize the given code, we can follow these steps:

- Use the ``Sequential`` API instead of the ``Model`` API for simplicity (since there are no branching or merging layers).
- Combine the data normalization and noise addition steps into a single loop to avoid unnecessary reshaping of the data.
- Use the ``binary_crossentropy`` loss directly in the ``evaluate`` method instead of using the ``predict`` method separately.

**FINAL PRACTICAL CODE:****1. PYTHON CODE**

```

import numpy as np
from sklearn.datasets import load_digits
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import MinMaxScaler
from tensorflow.keras.layers import Dense
from tensorflow.keras.models import Sequential
from tensorflow.keras import regularizers

# Load data
digits = load_digits()
X = digits.data # The input pixel values for each image
y = digits.target # The target (label) values for each image

# Normalize pixel values using MinMaxScaler
scaler = MinMaxScaler()
X = scaler.fit_transform(X)

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

# Define the RBM model using the Keras API
rbm = Sequential([
    Dense(128, input_shape=(64,), activation='relu',
    kernel_regularizer=regularizers.l2(0.01)),
    Dense(64, activation='relu', kernel_regularizer=regularizers.l2(0.01)),
    Dense(32, activation='relu', kernel_regularizer=regularizers.l2(0.01)),
    Dense(64, activation='relu', kernel_regularizer=regularizers.l2(0.01))
])

# Compile the RBM model
rbm.compile(optimizer='adam', loss='mean_squared_error')

# Train the RBM model
rbm.fit(X_train, X_train, epochs=5, batch_size=32)

# Extract features from RBM by computing the hidden layer activations for both
training and testing data
hidden_train = rbm.predict(X_train) # Hidden activations for training data
hidden_test = rbm.predict(X_test) # Hidden activations for test data

# Define the logistic regression classifier and hyperparameter search space
logreg = LogisticRegression(solver='lbfgs', max_iter=1000) # Classifier to
use for classification

```

```

params = {'C': np.logspace(-2, 2, 10)} # Range of hyperparameters to try

# Use grid search to find optimal hyperparameters and evaluate the performance
# of the model using cross-validation
clf = GridSearchCV(logreg, params, cv=5) # Perform a grid search over
hyperparameters using 5-fold cross-validation
clf.fit(hidden_train, y_train) # Fit the classifier to training data

# Evaluate the performance of the classifier on test data
y_pred = clf.predict(hidden_test) # Predict labels for test data using the
trained classifier
score = accuracy_score(y_test, y_pred) # Calculate the accuracy of predicted
labels compared to true labels
print('Accuracy:', score) # Print the accuracy score

```

## 2. CODE EXPLANATION

The code performs the following steps:

- *Import Necessary Libraries and Modules:*
  - ``numpy` as `np``: A library for numerical computing in Python.
  - ``load_digits` from `sklearn.datasets``: A function to load the digits dataset, which contains images of handwritten digits.
  - ``LogisticRegression` from `sklearn.linear_model``: A class for logistic regression, a classification algorithm.
  - ``accuracy_score` from `sklearn.metrics``: A function to calculate the accuracy of predicted labels.
  - ``train_test_split`, `GridSearchCV` from `sklearn.model_selection``: Functions/classes for splitting data into training and testing sets, and performing hyperparameter search using grid search with cross-validation, respectively.
  - ``MinMaxScaler` from `sklearn.preprocessing``: A class for normalizing pixel values.
  - ``Dense` from `tensorflow.keras.layers``: A class for fully connected neural network layers.
  - ``Sequential` from `tensorflow.keras.models``: A class for building sequential models in Keras.
  - ``regularizers` from `tensorflow.keras``: A module for applying regularization techniques in neural networks.
- Load the digits dataset using ``load_digits()``. The dataset consists of images of handwritten digits along with their corresponding target labels.
- Normalize the pixel values of the images using ``MinMaxScaler``, which scales the values between 0 and 1.
- Split the data into training and testing sets using ``train_test_split()``. The testing set will contain 20% of the data.
- Define an RBM (Restricted Boltzmann Machine) model using the Keras API. It is a type of unsupervised learning model with multiple layers of hidden units.

- Compile the RBM model with the Adam optimizer and mean squared error loss function.
- Train the RBM model by fitting it to the training data. The model is trained for 5 epochs with a batch size of 32.
- Extract features from the RBM by computing the hidden layer activations for both the training and testing data.
- Define a logistic regression classifier using `LogisticRegression`. This classifier will be used for classification.
- Define a grid containing hyperparameters to try for the logistic regression classifier. In this case, it defines a range of values for the inverse of the regularization strength (`C`) using a logarithmic scale.
- Use grid search cross-validation (`GridSearchCV`) to find the optimal hyperparameters for the logistic regression classifier. It performs a search over the predefined hyperparameter grid using 5-fold cross-validation.
- Fit the logistic regression classifier to the training data using the best hyperparameters found during the grid search.
- Use the trained classifier to predict labels for the hidden activations of the test data.
- Calculate the accuracy score by comparing the predicted labels with the true labels of the test data.
- Print the accuracy score.

Overall, the code trains an RBM model on the digits dataset and uses the extracted features as input to a logistic regression classifier. It then performs a grid search with cross-validation to find the best hyperparameters for the classifier and evaluates its performance on the test data. The final accuracy score is printed.

### 3. CODE COMPLETION

The code trains a Restricted Boltzmann Machine (RBM) for feature extraction and uses those features to train a logistic regression classifier. After training, it evaluates the performance of the classifier on test data.

### 4. CODE BUGS

Here are some problems with the given code:

- *Training the RBM Model using an Autoencoder-like Setup:* The code is training the RBM model using an autoencoder-like setup. Instead of predicting and comparing the input data ( $X_{train}$ ) with itself, it should be trained to reconstruct the input data based on the hidden activations.
- *Evaluation Metric for RBM Model Training:* The code is using the mean squared error (MSE) as the loss function for training the RBM model. However, MSE may not be the most suitable evaluation metric for generative models like RBMs. Consider using other metrics like reconstruction loss or free energy.

- *Missing Validation Set:* The code splits the data into training and testing sets, but it doesn't include a separate validation set for hyperparameter tuning. It's recommended to split the data into three sets: training, validation, and testing.
- *Missing Cross-Validation for RBM Model:* The code doesn't perform cross-validation for evaluating the performance of the RBM model. It's recommended to evaluate the model's performance using cross-validation to get a more robust estimate.

These are some problems with the given code that need to be addressed.

## 5. CODE OPTIMIZATION

To optimize the given code, we can follow these steps:

- Set the verbose parameter in the fit method to 0 to suppress the training progress output.
- Add the n\_jobs=-1 parameter to the GridSearchCV constructor to utilize all available CPU cores for parallel processing during hyperparameter search.

These optimizations make the code more concise and efficient.

## CONCLUSION:

In this practical experiment, we successfully implemented a Restricted Boltzmann Machine (RBM) and utilized it for dimensionality reduction on a given dataset. The RBM, as an unsupervised learning model, demonstrated its capability to effectively capture and represent meaningful features in the data while reducing its dimensionality.

Through the RBM training process, we observed the model's ability to learn the underlying probability distribution of the high-dimensional input data. By adjusting the weights using the contrastive divergence (CD) learning algorithm, the RBM minimized the energy function and produced a compressed representation of the data in its hidden layer.

The reduced representation obtained from the RBM's hidden layer activations showcased its potential in preserving essential information from the original data. Moreover, the reduced dimensionality facilitated improved data visualization, enabling us to gain insights into the dataset's structure and patterns more easily.

The experiment demonstrated that RBMs can be effective tools for dimensionality reduction tasks, particularly when dealing with high-dimensional data. The reduced representation offered a computationally more efficient solution for subsequent machine learning tasks, potentially reducing training time and memory requirements.

While the RBM showed promising results in capturing relevant features, it is essential to evaluate its performance and impact on specific downstream tasks. Depending on the application, the reduced representation may or may not be optimal for certain supervised learning tasks, such as classification. Further experimentation and evaluation on various datasets and tasks would provide deeper insights into the RBM's effectiveness as a dimensionality reduction technique.

In conclusion, the practical experiment highlighted the usefulness of RBMs for dimensionality reduction. By providing a lower-dimensional representation of the data, the RBM offers a viable approach to handling high-dimensional datasets efficiently. The successful implementation and analysis of the RBM's performance in this experiment open avenues for exploring its applications in various data analysis and machine learning scenarios, paving the way for further research and potential improvements in dimensionality reduction techniques.